




PM Role in an AI World: the session at a glance

120 minutes, the same in every region. One outcome: pilots people want to run on Monday.





How the breakout unfolds

One case per room, six steps, a checkpoint after each. Every room leaves with a Pilot Pack.

MINUTE	STEP	WHAT THE ROOM DOES	CHECKPOINT • WHAT MUST EXIST
00-05' 	1 Orient	Read the cold open and the case card aloud. Name the roles. The Builder opens Agent Builder.	Roles named; tool open, or paper mode chosen
05-15' 	2 Map the hours	Time each step of the job. Pick the biggest repetitive chunk and what those hours would buy.	Baseline in hours per PM per week, and its business impact
15-27' 	3 Design the agent	Brief the agent like a new colleague. Name what the PM decides, and the one way it could fail.	Instructions drafted, three tests written
27-42' 	4 Build & challenge two sub-teams in parallel	Builders create the agent and run the three tests. Challengers hunt the failure and write a guardrail.	Test results and review time recorded
42-50' 	5 Fix & decide	Add one rule per failed test. Choose one guardrail and one metric. Decide the level reached.	Final instructions, guardrail, level reached
50-60' 	6 Package & commit	Name the pilot owner and the first test. Rehearse the debrief card (90 seconds) once.	Pilot Pack submitted, debrief card ready
60' →	Summit Debrief opens the wrap-up	All rooms meet in the plenary. Each reads its card in 90 seconds; then one vote, three awards.	Monday-morning pilot chosen by vote

WHO IS IN EVERY ROOM

- 5-8 people, mixed experience
- Facilitator: runs process and clock
- Builder: builds the agent on screen

WHAT EVERY ROOM NEEDS

- A space close to the plenary
- Screen, network, capture link
- Printed case file and Pilot Pack

IF SOMETHING GOES WRONG

- Tool down over 3 min → “Be the agent” paper test
- Behind at 50' → cut detail, keep owner and guardrail
- Facilitator missing → the Scribe runs the room

Frequently asked questions

Left: before the summit. Right: on the day and after.



1 How are people assigned to rooms?

Upfront, by the regional lead, final on 26 Oct. The invitation (29 Oct) names each person's room and case file, C1-C6.

2 What must participants do before 5 November?

Read the story and their case file, map their own week, bring one idea. Done by Tue 3 Nov; nothing to prepare on the day.

3 Do participants need AI experience?

No. Each room needs one Builder with Agent Builder in Copilot Premium. Everyone else brings ideas and project experience.

4 What if a Builder doesn't have Agent Builder yet?

Request it at forms.cloud.microsoft/e/KifvXMztWp and confirm access by 3 Nov. In the room, the paper test is the fallback.

5 Who runs each room?

A facilitator named by the regional lead. They run the process and the clock, need no AI expertise, and are briefed on 23 Oct.

6 Which cases does my region run?

Every region runs C3 and C6, then adds C1, C4, C2, C5 in that order. Regions with six or more rooms run all six.

7 How many people per room?

5 to 8 with mixed experience. Above 10, split into two rooms on the same case. In rooms of five or fewer, people combine roles.

8 What does the site provide?

Per room: a space near the plenary, a screen, network, the capture link and printed backups. In plenary: the live scoreboard.

9 Which data do rooms use?

Only the KESTREL case files, which hold no real project data. No customer or confidential data goes into any AI tool.

10 What if the tool or network fails?

After 3 minutes, switch to "Be the agent": one person plays the agent on paper, so the room can still test its instructions.

11 What does each room deliver?

A Pilot Pack (agent instructions, hours back, guardrail, owner, first test) and a 90-second card for the Summit Debrief.

12 What happens after the summit?

Owners run their first test within two weeks. Packs are merged per case by 12 Nov; the 30-day readout follows in early December.