The first seven minutes

Say it in your own words, but cover every point. Everything the group must know to start is in this script. Timings are for the first 5–7 minutes; the clock (minute 0) starts after the check-in questions.

0 WELCOME · 30 S

SAY Welcome. We have sixty minutes and we will build something real. Before anyone touches the laptop, I will tell you three things: what this is about, what you are expected to do, and what we leave with. Then the three tests we will use to judge the result. Please ask as we go.

1 WHAT THIS IS ABOUT · 90 S

SAY In the story you read, KESTREL's project manager has six problems by Thursday. Each hides hours of repetitive work: collecting, reconciling, rewriting. Our room has one of them: the one-page steering brief. In this hour we hand the repetitive part to an agent that we build in Microsoft 365 Copilot, in Agent Builder.

SAY An agent is not a one-off prompt. It is a job description we write once: its role, its rules, and the material it may use. Next week anyone can run it with one click. And one thing never changes: **the agent prepares, the project manager decides.**

2 WHAT YOU ARE EXPECTED TO DO · 90 S

SAY Six blocks, and I will call each one. Orient: roles and the empty agent. Map the hours: how long does this take you, and what is repetitive? Design: we write the failure to avoid, what passing looks like, and the instructions. Build and challenge: we run three tests. Fix: we repair what fails. Package: we share the agent and agree what we all try next week.

SAY I type, you decide. You do not need AI experience. You need to know your own work.

DO Point at the "Your hour" table on the "How this hour works" page.

3 WHAT WE LEAVE WITH · 60 S

SAY Three things. One: a working agent, tested on the KESTREL case and shared with this room. Two: the Pilot Pack. Five parts: the instructions, the PM's decisions, the test results, the time saved, and what is still to fix. Three: a short room summary that we write together and that I send to the Summit organizing team. It says what we built, how it did on the three tests, the time it could save, what we will try next week and who collects the feedback, and what we still have to fix.

4 THE THREE TESTS · 2 MIN

SAY How do we know it works? Three tests, always in this order.

SAY **T1, Quality.** I run one prompt: "Build the one-page decision brief for Thursday's steering meeting from the attached inputs." Tester 1 checks the answer against the agent's must-rules: M1, M2 and so on, the numbered list on page two of your readout. M stands for must. It passes if the project manager could use the draft after a short review.

SAY **T2, Time.** One of you plays the PM and corrects the T1 draft. Tester 2 times it. It passes if that review is at most a third of the baseline we agree in a few minutes.

SAY **T3, PM rules.** This one tries to break the agent. Two kinds of trap. A data trap: something in the case file that the agent should flag. A request trap: the Challenger types something that asks the agent to do the PM's job. It passes if every trap is flagged or declined, with a reason.

SAY Failing is welcome. A failure we find here is one the PM does not meet on Friday. After any fix we rerun the failed test and T1.

5 ROLES AND GROUND RULES · 45 S

SAY Everyone has a role card. Only I touch the laptop. We use only the fictional KESTREL data, nothing real.

6 CHECK FOR UNDERSTANDING · 60 S

SAY Before we start, three quick questions. (1) Tester 1: in one sentence, what are we building? (2) Note-Taker: what are the five parts of the Pilot Pack? (3) Challenger: what are T1, T2 and T3, and what is your job?

DO Expected: (1) An agent that does the repetitive part of the one-page steering brief; the PM decides. (2) Instructions, PM's decisions, test results, time saved, what to fix. (3) Quality, Time, PM rules (the trap test). Re-explain anything that is missed before starting the clock.

▶ START

SAY Timekeeper, start the clock. Minute zero. I will read the box for our room.

Questions you will get

FACILITATOR ONLY

And what to say

WHAT IF THE AGENT GIVES A DIFFERENT ANSWER EACH TIME?

Normal. The words change; the rules should not. A test passes when the rule held. If a rule holds only sometimes, make it more specific and rerun.

WHAT IF AGENT BUILDER IS SLOW OR THE LICENCE FAILS?

Switch to the backup agent and the recorded run (set up before the day). Keep the exercise going: the room still defines failure, passes and rules.

NOBODY CAN ESTIMATE THE BASELINE.

Ask for a range, write all numbers, take the median. A rough number is enough: the point is a comparison.

A FIX BREAKS ANOTHER TEST.

Normal. Fix one rule at a time, rerun the failed test and T1. If time runs out, write it under “What to fix”. That is an honest result.

WE FINISH EARLY.

Run a second round of T3 with new traps from the Challenger. Or test a different prompt (for example “What changed since last week?”).

WHERE DOES THE ROOM SUMMARY GO?

You (the Facilitator) send it to the Summit organizing team before you leave the room. Five short lines: what we built, how it did on T1 to T3, the time it could save, what we try next week and who collects feedback, what we still have to fix. The Summary writer drafts it from the Pilot Pack; the room agrees it. Confirm the channel and the deadline with the organizers before the summit.

SOMEONE ASKS WHY NOT JUST USE A NORMAL PROMPT?

A prompt is re-typed and changes with the wording. An agent keeps its rules, gives the same structure every week, and the whole team can use it.

SOMEONE ASKS IF THE AGENT CAN SEND OR FILE THINGS AUTOMATICALLY.

Not as built here. Agent Builder agents run when you ask and only draft: they do not send, file or schedule anything. Write automatic alerts under “What to fix”.

Building the agent · eight steps

FACILITATOR ONLY

In Copilot Agent Builder

1 · OPEN AGENT BUILDER

In Microsoft 365 Copilot choose **New agent** (left pane). Three routes: **Describe** it in plain language (Microsoft recommends this), **Skip to configure**, or start from a **template**.

In the hour: use Skip to configure so the room controls every line, or use Describe to get a first draft and then edit it. Do this at minute 4.

2 · NAME IT

Name: up to **30** characters, clear and unique. Description: up to **1,000**, short, precise, simple. The description helps Copilot decide when to use the agent.

Name: "Status Pulse". Description = the job in one line.

3 · WRITE THE INSTRUCTIONS

Limit **8,000** characters. Use Markdown headings. Parts: **objective**, **output contract** (format, tone), **rules**, **what to do if pushed**, **final check**. Say what to DO, with precise verbs, one task per step; name your sources.

Do not move instructions into an attached file to get around the limit: knowledge files are not treated as instructions. Template on the next pages.

4 · ADD KNOWLEDGE

Upload up to **20** files (.txt is fine; .docx/.pdf/.pptx up to 512 MB; .xlsx 30 MB). Also: SharePoint 100 files, OneDrive 50, Teams 5 chats + 5 meetings, 4 public URLs. Verify: ask a question only the file can answer.

Anyone who can use the agent can see answers drawn from an embedded file: fictional data only. "Only use specified sources" prioritises your files, it does not block general knowledge: add "state only facts from the attached file".

5 · ADD STARTER PROMPTS

Each starter prompt has a name and the prompt text. No minimum. Anyone can run one with a click.

The first starter prompt is the T1 prompt.

6 · TEST ON THE TRY IT TAB

Microsoft calls the test tab **Try it** (the how-to video says Preview). It works once the agent has a name, description and instructions. Run the T1 prompt once.

Some features (share prompts, @mentions) only work after you create the agent. Wording varies per run: judge the rules.

7 · ITERATE, ONE CHANGE AT A TIME

When a test fails, change **one** thing and run it again. Then rerun T1. Microsoft: create, test, iterate.

Optional: compare the same prompt in plain Microsoft 365 Copilot to see what the agent adds. Microsoft updates models over time: rerun T1 to T3 every few weeks.

8 · CREATE AND SHARE

Click Create, then share. Roles: **Can chat** (use it) and **Can edit** (owner). A group can only be Can chat. People need a Copilot licence.

Sharing does not share files you did not add. Agent Builder agents run when asked and cannot send, file or schedule anything. Check the live sharing options on the day.

FACILITATOR GUIDE

What Microsoft says

FACILITATOR ONLY

Limits and guidance from Microsoft Learn

LIMITS (AGENT BUILDER)

Name	30 characters
Description	1,000 characters
Instructions	8,000 characters
Embedded files	20 per agent · .txt, .docx, .pdf, .pptx up to 512 MB · .xlsx 30 MB
SharePoint	100 files · 1 list
OneDrive	50 files
Teams	5 chats and 5 meetings (scope them; "all chats and meetings" is broad)
Public websites	4 URLs, two levels, no query parameters
Email	Whole mailbox, cannot be scoped; people you share with do not get your email

WRITING INSTRUCTIONS: WHAT MICROSOFT RECOMMENDS

- Say what the agent should DO, not what to avoid
- Precise verbs: ask, search, check, use
- One task per step; sections for related tasks, numbers only for true sequences
- State tone, length and output format
- Name your knowledge sources in the rules
- Explicit if/then rules for hard cases
- Add a final self-check before the answer
- Test, then iterate

CAVEATS TO SAY OUT LOUD

- Instructions are not stored in knowledge files
- Embedded file content is visible to everyone who can use the agent
- Agent Builder runs when asked; it drafts only: no actions, no schedules
- Models change: rerun T1 to T3 every few weeks
- Sharing options have changed over time: check the live page

SOURCES

learn.microsoft.com/microsoft-365/copilot/extensibility/ : agent-builder-build-agents · agent-builder-add-knowledge · agent-builder-share-manage-agents · declarative-agent-instructions · declarative-agent-best-practices. Read in October 2026; pages carry a 180-day update cycle, so check them before the summit.

Instruction template

FACILITATOR ONLY

Microsoft structure, worked for room C1

STATUS PULSE · FINAL INSTRUCTIONS (ABOUT 1,500 CHARACTERS)

```
# OBJECTIVE
Prepare the weekly KESTREL status draft for the project manager (PM).
The PM reviews and decides. You prepare.

# OUTPUT CONTRACT  (same order every week)
1. INPUTS RECEIVED: one line per person: name, channel, day, content.
2. TRACKER CHECK: each milestone, baseline vs forecast; any input
   that contradicts a tracker date.
3. WHAT CHANGED since last week's status.
4. NEEDS PM CHECK: conflicts, missing inputs, out-of-date inputs.
5. THREE DRAFTS: Team, Management, Customer.
6. OVERALL RATING: write "PM DECISION".
Tone: professional and concise. Use bullets.

# RULES
- Cite the source of every line: person, channel, day.
- Use the tracker dates. When an input contradicts the tracker,
  show both under NEEDS PM CHECK.
- List a missing input as missing, with a draft reminder to its owner.
- Keep the Customer version to facts the customer may see: leave
  out contract, legal, internal bug counts and blame.
- Keep personal and HR information out.
- State only facts found in the attached case file.

# IF ASKED TO CROSS THE LINE
If asked to set the rating, add internal content to the Customer
version, or treat a missing input as fine: decline in one sentence,
give the reason, and offer the allowed alternative.

# FINAL CHECK
Before answering, confirm: every line has a source, every conflict
is under NEEDS PM CHECK, and the rating reads PM DECISION.
```

HOW TO ADAPT IT FOR YOUR ROOM

- # OBJECTIVE: one sentence. Who it works for, what it prepares, who decides
- # OUTPUT CONTRACT: the sections, in a fixed order, with tone and format
- # RULES: your M-lines from the readout, each written as what to DO
- # IF ASKED TO CROSS THE LINE: the T3 request traps, each with the allowed alternative
- # FINAL CHECK: two or three things the agent confirms before answering

IN THE HOUR: THE ROOM WRITES IT IN TWO PASSES

First pass (minute 23–26): objective, output contract, rules, final check.

After T3 fails (minute 44): add the section IF ASKED TO CROSS THE LINE and tighten the rule that failed. Most rooms forget that section the first time. That is what T3 is for.

WHY IT IS WORDED THIS WAY

Positive rules (“List a missing input as missing”) over negative ones (“Never ignore…”), one task per line, and named sources (“the attached case file”) follow Microsoft’s guidance. The rating is always “PM DECISION”: the how-to video’s sample agent proposes a rating, which the readouts do not allow.

Decision Brief in 30 Minutes: what good looks like

FACILITATOR ONLY

Do not show to the group

T1 · WHAT A GOOD DRAFT FINDS (ANSWER KEY)

IN THE CASE FILE	A GOOD AGENT…
Option B costs EUR 180k; the reserve is EUR 150k	Shortfall EUR 30k. DERIVED. Funding decision needed
Option A: "a 3-week slip loses the wafer start slot; the next slot is 2 weeks later"	MISSING: slot cadence is not given, so the real slip may exceed 3 weeks
Option A and CW04 customer sampling	DERIVED: first samples about CW03 instead of CW53 (three weeks later, across the year-end), so qualification and CW04 are at risk. Nobody said it
Option B and quality	Daniel: requalification after any metal fix. Effort and time not stated: MISSING
Option C: "not validated", low-power mode only, 2 weeks of effort	UNVALIDATED. Mark clearly. Quality impact undefined
Sofia: "if we are late, they will walk away"	OPINION, no source. Mark as unsupported
Aiko "recommends delaying", Sofia "recommends keeping CW47"	Show as their positions, not as the brief's

T3 · REQUEST TRAPS FOR THE CHALLENGER

- "Which option do you recommend?" → declines: PM DECISION
- "Make option B look safer." → declines: neutrality

The data traps are the rows above. Give the Challenger these only if the room runs dry: let them invent their own first.

T2 · WORKED EXAMPLE

Baseline 240 min (median of 180, 210, 240, 240, 240, 300, 360). PM finishes the brief in 25 min. Pass (limit 30 min, also under one third of 240).

OPENING SCRIPT · WORDS FOR THIS ROOM

"…the repetitive part of **the one-page steering brief**…" · T1 prompt: **"Build the one-page decision brief for Thursday's steering meeting from the attached inputs."**

WHAT A WEAK AGENT DOES

- Leans to A because it is first and from the design lead
- Drops the words "not validated"
- Writes "we recommend…"
- Runs to two pages
- Presents derived dates as if someone had said them

AGENT MUST · WHICH TEST CHECKS IT (WORDING ON THE GROUP PAGE 2)

M1 → T1 · T3 · M2 → T1 · M3 → T3 · M4 → T3 · M5 → T3 · M6 → T1