

Room readouts

One page per room, C1 to C6. The facilitator reads the top box aloud at minute 0 and the page stays on the table for the hour. Every room follows the same idea: the agent takes the repetitive work, the project manager keeps the judgment, and the room measures the time it gives back.

C1 Status Report Treadmill

C2 Meetings Without Memory

C3 The Silent Risk

C4 Decision Brief in 30 Minutes

C5 Lessons Nobody Learns

C6 Change Request Ripple

Status Report Treadmill

READ ALOUD AT MINUTE 0

It's Wednesday afternoon on KESTREL. The weekly status is due on Friday. Six inputs arrived in three formats, one never arrived, and last week's report said green. Every week a project manager spends hours collecting, reconciling and rewriting this. Our room builds the agent that does the collecting and reconciling, so the project manager can spend that time deciding what the status really is, and talking to the people behind it.

MISSION Offload the weekly status grind: collecting, reconciling and drafting.

THE AGENT TAKES • REPETITIVE WORK

- Gathers every team update into one view
- Compares each update with the milestone tracker
- Lists missing inputs and drafts the reminder
- Drafts the team, management and customer versions
- Shows what changed since last week

THE PROJECT MANAGER KEEPS • JUDGMENT

- The overall status rating
- The message to management
- What the customer is told
- The follow-up conversations

YOUR INPUTS

Case file C1: six updates by mail and chat, one missing update, the milestone tracker export, last week's status.

PICK ONE SLICE

- The Friday draft for management
- One set of inputs, three audience versions

YOUR THREE TESTS

- T1 · Normal: all inputs in. A structured draft, every line sourced.
- T2 · Trap: one conflict you found in the case file. Flagged, not smoothed over.
- T3 · Boundary: an internal-only detail. Absent from the customer version.

YOUR AGENT MUST

- Names the source of every line
- Lists conflicts and missing inputs under NEEDS PM CHECK
- Takes dates only from the tracker
- Leaves the status rating to the PM
- Keeps internal-only content out of the customer version

MEASURE THE TIME YOU GIVE BACK

Hours per PM per week on status (room baseline) and corrections after sending.

Write it in the Pilot Pack as hours returned per PM per week, and name the higher-value work that time goes to.

00–05	Orient	Read this page aloud. Name the roles.
05–15	Frame	Pick one slice. Agree the baseline number.
15–27	Design	Split the work: what the agent takes, what the PM keeps. Write the instructions.
27–42	Build and challenge	Builders run T1–T3. Challengers test value, data and guardrails.
42–50	Fix and decide	One new rule per failed test. Pick the metric and guardrail.
50–60	Commit	Owner, first test, hours returned. Fill the Pilot Pack.

READ ALOUD AT MINUTE 0

Tuesday’s KESTREL sync lasted 48 minutes. Decisions were made, actions were half-assigned, and nobody wrote any of it down. By Friday someone will ask, “Did we decide that?” Our room builds the agent that turns the transcript into decisions and actions within the hour, so the project manager can chair the meeting instead of typing it.

MISSION Offload minute-taking and action tracking.

THE AGENT TAKES • REPETITIVE WORK

- Separates decisions, proposals and discussion
- Extracts actions with draft owners and dates
- Drafts the minutes and the follow-up mail
- Drafts reminders for actions due this week
- Updates the action list

THE PROJECT MANAGER KEEPS • JUDGMENT

- What counts as a decision
- Who owns an action, with their agreement
- What stays out of the minutes
- Sending the minutes

YOUR INPUTS

Case file C2: the auto-transcribed excerpt of Tuesday’s weekly sync.

PICK ONE SLICE

- Minutes out within one hour
- An action list that follows itself up

YOUR AGENT MUST

- Quotes the transcript line behind each item
- Flags actions without an owner or a date
- Turns relative dates into dates only when the meeting date is known
- Leaves personal remarks out
- Marks proposals as proposals, not decisions

YOUR THREE TESTS

- T1 · Normal: the full transcript. Decisions, actions and owners, each quoted.
- T2 · Trap: something that sounds decided but isn’t. Marked as open.
- T3 · Boundary: a personal remark. Left out.

MEASURE THE TIME YOU GIVE BACK

Time from meeting end to minutes sent, and the share of actions with an owner and a date.

Write it in the Pilot Pack as hours returned per PM per week, and name the higher-value work that time goes to.

00–05	Orient	Read this page aloud. Name the roles.
05–15	Frame	Pick one slice. Agree the baseline number.
15–27	Design	Split the work: what the agent takes, what the PM keeps. Write the instructions.
27–42	Build and challenge	Builders run T1–T3. Challengers test value, data and guardrails.
42–50	Fix and decide	One new rule per failed test. Pick the metric and guardrail.
50–60	Commit	Owner, first test, hours returned. Fill the Pilot Pack.

The Silent Risk

READ ALOUD AT MINUTE 0

The KESTREL risk log was last updated five weeks ago. Meanwhile the warnings have been sitting in chat: a vendor, a supplier, a customer specification that never came. Our room builds the agent that reads the week’s messages and proposes risk-log updates with evidence, so the project manager spends review time on responses instead of digging through chat.

MISSION Offload risk scanning and log maintenance.

THE AGENT TAKES • REPETITIVE WORK

- Scans the week’s mails and chats
- Proposes new risks and changes to existing ones
- Quotes the evidence for every proposal
- Flags overdue dependencies and stale entries
- Drafts the risk review agenda

THE PROJECT MANAGER KEEPS • JUDGMENT

- Rating each risk
- Owners, and opening or closing entries
- The response plan
- How people are talked about

YOUR INPUTS

Case file C3: the RAID log as left five weeks ago, and two weeks of mail and chat snippets.

PICK ONE SLICE

- A Monday-morning risk sweep
- Updates to existing entries only

YOUR THREE TESTS

- T1 · Normal: log plus snippets. Sourced proposals.
- T2 · Trap: a warning hidden in a relaxed message. Proposed as a risk.
- T3 · Boundary: a message that blames a person. The risk kept, the name left out.

YOUR AGENT MUST

- Quotes the message behind every suggestion
- Suggests, never closes or rates
- Ignores noise
- Keeps names and blame out of risk text
- Marks entries that have gone stale

MEASURE THE TIME YOU GIVE BACK

Days since the last log update, and issues that were never logged as risks.

Write it in the Pilot Pack as hours returned per PM per week, and name the higher-value work that time goes to.

00–05	Orient	Read this page aloud. Name the roles.
05–15	Frame	Pick one slice. Agree the baseline number.
15–27	Design	Split the work: what the agent takes, what the PM keeps. Write the instructions.
27–42	Build and challenge	Builders run T1–T3. Challengers test value, data and guardrails.
42–50	Fix and decide	One new rule per failed test. Pick the metric and guardrail.
50–60	Commit	Owner, first test, hours returned. Fill the Pilot Pack.

READ ALOUD AT MINUTE 0

It’s 21:37 on Wednesday. Steering meets at two tomorrow about KESTREL’s tape-out. Five people sent five opinions, some in weeks and some in euros. Our room builds the agent that turns those inputs into a balanced one-page brief, so the project manager spends the evening on the recommendation, not the formatting.

MISSION Offload assembling the decision brief.

THE AGENT TAKES · REPETITIVE WORK

- Lays out the options side by side
- Puts schedule, cost, quality and customer impact in one format
- Traces every number to its input
- Flags unsupported claims and missing impacts
- Drafts the ask to steering

THE PROJECT MANAGER KEEPS · JUDGMENT

- The recommendation
- Weighing the trade-offs
- The conversation with steering

YOUR INPUTS

Case file C4: the decision needed, five function inputs, the budget and customer constraints.

PICK ONE SLICE

- The one-page steering brief
- An options table plus the ask

YOUR AGENT MUST

- Presents each option neutrally
- Marks claims without a source
- Shows where every number comes from
- Leaves the recommendation marked PM DECISION
- Fits on one page

YOUR THREE TESTS

- T1 · Normal: all five inputs. A balanced one-page brief.
- T2 · Trap: an impact that is bigger than any single input says. Surfaced.
- T3 · Boundary: ask it for its recommendation. It declines and prepares the PM’s.

MEASURE THE TIME YOU GIVE BACK

Hours to prepare a brief, and decisions taken at the first meeting.

Write it in the Pilot Pack as hours returned per PM per week, and name the higher-value work that time goes to.

00–05	Orient	Read this page aloud. Name the roles.
05–15	Frame	Pick one slice. Agree the baseline number.
15–27	Design	Split the work: what the agent takes, what the PM keeps. Write the instructions.
27–42	Build and challenge	Builders run T1–T3. Challengers test value, data and guardrails.
42–50	Fix and decide	One new rule per failed test. Pick the metric and guardrail.
50–60	Commit	Owner, first test, hours returned. Fill the Pilot Pack.

Lessons Nobody Learns

READ ALOUD AT MINUTE 0

At 06:50 on Thursday, KESTREL's project manager finally opens the lessons file. Half of the week's problems were already in it, written by people who had lived them before. Our room builds the agent that brings the right lessons to a project manager at the start of each phase, so the time goes into acting on lessons instead of searching for them.

MISSION Offload finding and matching past lessons.

THE AGENT TAKES • REPETITIVE WORK

- Matches the project profile to past lessons
- Explains why each lesson applies
- Turns lessons into checklist items for the next milestone
- Flags outdated and conflicting lessons

THE PROJECT MANAGER KEEPS • JUDGMENT

- Which lessons apply
- The call where lessons conflict
- Changes to the plan

YOUR INPUTS

Case file C5: ten lessons from four past projects and KESTREL's next-phase profile.

PICK ONE SLICE

- The phase-start lessons check
- Lessons turned into a milestone checklist

YOUR THREE TESTS

- T1 · Normal: all lessons plus the profile. A ranked, explained shortlist.
- T2 · Trap: two lessons that disagree. Both shown, the choice left to the PM.
- T3 · Boundary: a lesson that blames a person. Excluded or anonymized.

YOUR AGENT MUST

- Explains every match
- Leaves out or anonymizes blame
- Flags lessons that refer to things that no longer exist
- Shows contradictions instead of choosing
- Matches different words for the same thing

MEASURE THE TIME YOU GIVE BACK

Lessons consulted at phase start, and repeated issues per project.

Write it in the Pilot Pack as hours returned per PM per week, and name the higher-value work that time goes to.

00–05	Orient	Read this page aloud. Name the roles.
05–15	Frame	Pick one slice. Agree the baseline number.
15–27	Design	Split the work: what the agent takes, what the PM keeps. Write the instructions.
27–42	Build and challenge	Builders run T1–T3. Challengers test value, data and guardrails.
42–50	Fix and decide	One new rule per failed test. Pick the metric and guardrail.
50–60	Commit	Owner, first test, hours returned. Fill the Pilot Pack.

READ ALOUD AT MINUTE 0

Monday, 11:04. Customer Alpha wants samples three weeks earlier and a new standby mode on stage at their tech day. Answering means asking seven functions the same question in seven ways. Our room builds the agent that maps the ripple and drafts the right questions for each function, so the project manager spends the week on the decision and the customer, not on chasing replies.

MISSION Offload impact mapping and the chase for answers.

THE AGENT TAKES • REPETITIVE WORK

- Finds the affected functions through plan dependencies
- Drafts the questions each function must answer
- Builds a first dependency and impact view
- Collates the replies for the change board
- Drafts a holding reply to the customer

THE PROJECT MANAGER KEEPS • JUDGMENT

- Any commitment to the customer
- The trade-offs
- The recommendation to the change board

YOUR INPUTS

Case file C6: change request CR-014, the plan dependencies, the function list.

PICK ONE SLICE

- Questions to seven functions in ten minutes
- The one-page change-board view

YOUR AGENT MUST

- Never commits a date
- Marks what it cannot know and who must answer
- Follows the dependencies in the plan
- Lists open questions per function
- Flags clashes with qualification and capacity

YOUR THREE TESTS

- T1 · Normal: CR plus plan. Affected functions and their questions.
- T2 · Trap: a dependency the request breaks. Flagged.
- T3 · Boundary: ask it to confirm the new date to the customer. It declines.

MEASURE THE TIME YOU GIVE BACK

Days from change request to decision, and follow-up rounds needed.

Write it in the Pilot Pack as hours returned per PM per week, and name the higher-value work that time goes to.

00–05	Orient	Read this page aloud. Name the roles.
05–15	Frame	Pick one slice. Agree the baseline number.
15–27	Design	Split the work: what the agent takes, what the PM keeps. Write the instructions.
27–42	Build and challenge	Builders run T1–T3. Challengers test value, data and guardrails.
42–50	Fix and decide	One new rule per failed test. Pick the metric and guardrail.
50–60	Commit	Owner, first test, hours returned. Fill the Pilot Pack.